

An Integer Programming Approach to Solving Rummikub Configurations

M. Zermeno
School of Physics and Mathematics
UANL

May 14, 2025

Abstract

Rummikub is a popular tile-based game that involves both combinatorial optimization and strategic decision-making. This paper presents a formal investigation into the game mechanics, identifies the core computational challenges, and introduces an integer programming (IP) model designed to find optimal moves for a given player state. We analyze the problem space, define constraints based on legal game rules, and propose an IP formulation to maximize score per turn. Giving a future insight into applying possible Counting strategies to optimize decision making when given multiple players' states.

1 Introduction

Rummikub is a widely played tile-based game that combines elements of rummy and mahjong. Players aim to be the first to empty their racks by forming sets and runs using numbered and colored tiles. While the rules are simple, the combinatorial complexity that emerges as the game progresses makes decision-making increasingly difficult. A player may have many possible ways to arrange their tiles on the board, particularly when they are allowed to manipulate existing sets. Deciding on the best possible move—i.e., the one that maximizes the number or value of tiles placed on the table—is a non-trivial problem.

In this research, we investigate the use of Integer Linear Programming (ILP) to solve the Rummikub move optimization problem. Our goal is to develop a mathematical model that, given a current board state and a player's hand, determines the optimal set of tile moves that obey all Rummikub rules while maximizing a strategic objective, such as the number or value of tiles placed. This approach is based on the successful work of Den Hertog and Hulshof (2006), who showed that Rummikub configurations can be effectively modeled and solved using ILP within seconds, even in nontrivial game states.

1.1 Why use Integer Linear Programming?

Integer Linear Programming is a branch of mathematical optimization used for solving decision problems involving discrete variables and linear relationships. In ILP, the goal is to maximize or minimize a linear objective function (like total tile value placed), subject to a set of linear equality or inequality constraints (like the legal formation of sets).

In Rummikub, each decision—whether to play a tile, form a group or a run, or manipulate the board—can be encoded as a binary or integer variable (e.g., "1" if the tile is used, "0" if not). The constraints ensure that moves are legal according to the game rules (e.g., sets must contain valid combinations of numbers and colors, no tile can be used more than once, jokers are used properly, etc.).

Using ILP offers several advantages:

- **Optimality:** It guarantees the best move given the current information.
- **Flexibility:** Complex rules (like joker usage and minimizing board changes) can be incorporated as constraints or secondary objectives.
- **Speed:** With modern solvers, even complex configurations can be solved in under a second.

This is especially valuable in Rummikub because brute-force enumeration of all legal move combinations becomes computationally infeasible as the number of tiles increases. We will further explore the brute-force enumeration search space, when discussing the ILP model.

2 Background and Related Work

2.1 Rummikub Game Rules Overview

Rummikub is a tile-based game for 2 to 4 players (or up to 6 with extended sets). The goal is to be the first to eliminate all tiles from one's rack by forming valid sets according to specific rules. Below, we summarize the essential components and gameplay mechanics relevant to our mathematical modeling.

2.2 Objective

The aim of the game is to place all the tiles from one's rack onto the table by forming valid sets: **groups** or **runs**.

2.3 Tile Set

The complete game includes:

- 106 tiles: two sets of tiles numbered 1 to 13 in four colors (red, blue, black, and orange), plus 2 jokers.
- Each player starts with 14 tiles on a personal rack.
- Remaining tiles form a face-down pool.

2.4 Valid Sets

- **Group:** Three or four tiles of the same number in different colors.
- **Run:** Three or more consecutive numbers in the same color.

2.5 Setup

- Players draw one tile each; the highest goes first.
- All tiles are then shuffled and stacked in piles (optional for convenience).
- Each player draws 14 tiles for their rack.

2.6 Initial Meld

- A player's first move must consist of one or more valid sets with a combined value of at least 30 points.
- Tiles used must come solely from the player's rack.
- Jokers may be used and are valued as the tile they represent.
- If unable or unwilling to make the initial meld, the player draws a tile from the pool and ends their turn.

2.7 Gameplay

- Turns proceed clockwise.
- After the initial meld, players may:
 - Add tiles to existing sets on the table.
 - Rearrange existing sets (manipulation) to create new valid sets.
- If no valid move is possible, the player must draw a tile and end their turn.
- Tiles drawn cannot be played until the player's next turn.

2.8 Manipulation Rules

Players may manipulate sets on the table under the following conditions:

- At the end of the turn, all tiles on the table must be in valid sets.
- No tiles may remain loose or isolated.
- Jokers may be replaced if the player can substitute them with the correct tile.
- The joker must then be immediately reused to form a new set in the same turn.

2.9 The Joker

- Acts as a substitute for any tile in a valid set.
- Can be reclaimed if replaced by the tile it represents.
- Must be played immediately in a new set using at least one tile from the player's rack.
- Cannot be used before the player's initial meld is completed.

2.10 Winning and Scoring

- A game ends when a player empties their rack.
- Other players total the value of the tiles remaining on their racks; this becomes their negative score.
- The winner scores the sum of the other players' negative scores as a positive amount.
- The joker carries a penalty of 30 points if left on a rack.
- If no player can make a move and the pool is empty, the player with the lowest rack total wins.

2.11 Computational Complexity

Previous research suggests that determining a valid move is NP-hard due to the permutations and combinations of tile arrangements. de Jong and Uiterwijk [2014].

2.12 Related AI Models

Discuss existing models or papers on game AI, such as those used in Sudoku, Scrabble, or other combinatorial games.

3 Problem Definition and Set Enumeration

3.1 Since the Problem We Are Solving

The core decision-making challenge in a game of Rummikub is determining the best possible move during a player's turn. Specifically, the question is:

What is the largest number (or total value) of tiles that a player can legally place on the table in a single turn, either by forming new sets or by manipulating existing sets, in accordance with all Rummikub rules?

This is a nontrivial problem due to the combinatorial explosion of possible tile groupings and manipulations. A single rack of 14 tiles, combined with the dynamic state of the board, can result in thousands of possible legal moves. Attempting to evaluate all possible combinations by hand or brute-force computation is computationally infeasible. Thus, the use of **Integer Linear Programming (ILP)** offers a structured way to model these constraints and systematically identify optimal solutions.

To build such a model, it is essential to first define the full universe of legal **Rummikub sets** (i.e., all combinations of tiles that can be legally played), which is where the number **1174** plays a critical role.

3.2 Enumeration of All Valid Sets

In Rummikub, a **set** is a valid grouping of tiles that satisfies one of the following two rules:

- **Group:** Three or four tiles of the same number in different colors.

- **Run:** Three to five¹ consecutive numbers in the same color.

The authors of the original model compute the total number of valid sets that can occur in the game—whether made purely from natural tiles or involving jokers. This total comes out to exactly **1174 sets**. Below, we explain how this number is derived.

1. Sets Without Jokers

- **Runs (same color, consecutive numbers):**
 - Each color (4 in total) allows for:
 - * 11 runs of 3 numbers (e.g., 1–3, 2–4, ..., 11–13)
 - * 10 runs of 4 numbers
 - * 9 runs of 5 numbers
 - Total runs without jokers: $44 + 40 + 36 = 120$
- **Groups (same number, different colors):**
 - For each number from 1 to 13:
 - * Choose 3 out of 4 colors: $\binom{4}{3} = 4$ combinations $\rightarrow 13 \times 4 = 52$
 - * Or all 4 colors $\rightarrow 13$ combinations
 - Total groups without jokers: $52 + 13 = 65$

2. Sets With Jokers

Jokers can substitute for any missing tile in a group or run, significantly increasing the number of possible valid sets. The authors enumerated all legal joker-containing combinations by hand and through exhaustive generation, ensuring no illegal duplications.

- **With one joker:**
 - 92 valid 3-tile runs
 - 124 valid 4-tile runs
 - 148 valid 5-tile runs
 - 78 valid 3-tile groups
 - 52 valid 4-tile groups
 - Total: $92 + 124 + 148 + 78 + 52 = 494$
- **With two jokers:**
 - 52 valid 3-tile runs
 - 132 valid 4-tile runs
 - 233 valid 5-tile runs
 - 78 valid 4-tile groups
 - Total: $52 + 132 + 233 + 78 = 495$

¹Note that a run with six or more consecutive numbers can be divided into runs of length three, four or five, and so we do not need to consider them separately

3. Final Total

The total number of valid Rummikub sets across all cases is:

$$\text{No jokers: } 185 \quad + \quad \text{1 joker: } 494 \quad + \quad \text{2 jokers: } 495 \quad = \boxed{1174 \text{ sets}}$$

These 1174 sets become the foundation of the model: the optimization process selects which sets (from this pool) to use to maximize tile usage on a given turn.

3.3 Why This Enumeration Matters

Each set is assigned an index ($j \in \{1, \dots, 1174\}$) and is represented in the ILP model using binary or integer variables. By precomputing the entire universe of legal sets, the model transforms the abstract problem of Rummikub tile play into a structured selection problem with known components, enabling the solver to:

- Decide which sets to form (and how many times),
- Ensure all tiles are used legally,
- Maximize the number (or value) of tiles placed,
- Optionally minimize unnecessary changes to the table.

4 Model Details

4.1 Overview

To determine the optimal move in a given Rummikub configuration, Den Hertog and Hulshof proposed an Integer Linear Programming (ILP) model. The model computes the best possible placement of tiles from a player's rack onto the table, optionally involving manipulation of existing sets. The optimization goal is to either:

- Maximize the **number of tiles** placed, or
- Maximize the **value of tiles** placed (i.e., sum of tile numbers).

4.2 Model Structure

The model defines the following sets, parameters, and variables:

Indices:

- $i \in I = \{1, \dots, 53\}$: index for tile types (including the joker).
- $j \in J = \{1, \dots, 1174\}$: index for all possible valid sets (runs and groups, with or without jokers).

Parameters:

- s_{ij} : equals 1 if tile i is part of set j ; 0 otherwise.
- t_i : number of times tile i is currently on the table.
- r_i : number of times tile i is available in the player's rack.
- v_i : value of tile i (equal to its number; joker can have custom value).

Variables:

- $x_j \in \{0, 1, 2\}$: number of times set j is formed in the new solution.
- $y_i \in \{0, 1, 2\}$: number of times tile i is played from the player's rack.

4.3 Objective Function

The primary objective is to maximize the number of tiles placed from the rack:

$$\max \sum_{i \in I} y_i$$

Alternatively, the model can maximize the total value of those tiles:

$$\max \sum_{i \in I} v_i \cdot y_i$$

4.4 Constraints**Tile Conservation Constraint:**

$$\sum_{j \in J} s_{ij} x_j = t_i + y_i \quad \forall i \in I$$

This ensures that each tile used in the final solution either:

- Was already present on the table, or
- Was placed from the player's rack.

Rack Availability Constraint:

$$y_i \leq r_i \quad \forall i \in I$$

A tile cannot be played more times than it is available in the player's rack.

Variable Domains:

$$x_j \in \{0, 1, 2\}, \quad y_i \in \{0, 1, 2\}$$

Each set or tile can appear at most twice, because there are at most two of each tile in the game.

4.5 Feasibility

The model is always feasible, since setting $y_i = 0$ for all i results in the current state of the board being preserved. This baseline corresponds to making no move at all, yielding an objective value of zero.

5 Implementation and Experiments

5.1 Tools Used

Model implemented using Python and solved with Google OR-Tools

5.2 Test Cases

Present example game states and corresponding optimal plays derived from the model.



Figure 1: Example board configuration solved by the model.

5.3 Results

Show performance in terms of solution time, and move quality.

6 Discussion

6.1 Limitations

- Model scalability as board complexity increases
- Handling dynamic player interaction

6.2 Future Work

- Integration with other AI agents.
- Use of reinforcement learning or hybrid approaches.
- Implementation of Counting strategies to create a complex probabilistic model.

7 Conclusion

We have proposed an integer programming model capable of identifying optimal Rummikub moves from a static game state. This framework serves as a step toward developing intelligent agents capable of playing competitively.

References

- E. de Jong and J. Uiterwijk. Rummikub is np-complete. *ICGA Journal*, 37(3):151–157, 2014.